NES

David Kane

2/26/2020

Model Results

Here is a nicely printed summary of my model results, showing the relationship between party ID, on a 1 through 7 scale, and ideology, also on a 1 through 7 scale. Although these are ordered categories, I treat both variables as continuous.

```
            Median MAD_SD
(Intercept) 0.97   0.05
real_ideo   0.66   0.01

Auxiliary parameter(s):
      Median MAD_SD
sigma 1.90   0.01
```

It isn't the best thing I could imagine. I really want to figure out a way to add a caption. But it will do for now.

Instead of just printing the simple object, we could print its summary.

```
Model Info:
 function:     stan_glm
 family:       gaussian [identity]
 formula:      partyid7 ~ real_ideo
 algorithm:    sampling
 sample:       4000 (posterior sample size)
 priors:       see help('prior_summary')
 observations: 16590
 predictors:   2

Estimates:
              mean   sd   10%   50%   90%
(Intercept) 0.97   0.05 0.90  0.97  1.03
real_ideo   0.66   0.01 0.65  0.66  0.68
sigma       1.90   0.01 1.88  1.90  1.91

Fit Diagnostics:
           mean   sd   10%   50%   90%
mean_PPD 3.79   0.02 3.76  3.79  3.82

The mean_ppd is the sample average posterior predictive distribution of the outcome variable (for detail

MCMC diagnostics
              mcse Rhat n_eff
```

```
(Intercept)   0.00 1.00 3799
real_ideo     0.00 1.00 4074
sigma         0.00 1.00 3651
mean_PPD      0.00 1.00 3470
log-posterior 0.03 1.00 1769

For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective sample si
```

But that adds a bunch of junk, including run-off-the-page text. Maybe there are options in something like `print.stanmvreg()` which might be helpful.